Assignment8_Seshadri

Sri Seshadri

8/12/2017

Assignment

In this assignment we will learn how to perform an exploratory data analysis for a clustering problem, fit a hierarchical cluster analysis, fit a k-means cluster analysis, how to integrate principal components analysis and cluster analysis, how to use cluster analysis as a predictive model, and how to make a variety of R graphics applicable to cluster analysis and multivariate analysis in general.

Part 1- The data

Let's begin by reading in the data.

```
my.data <- read.csv("European_Employment.csv", header = T)
levels(my.data$Group)
```

```
## [1] "Eastern" "EFTA"    "EU"      "Other"
```

```
pairs(my.data[, -2], col = my.data$Group)
```

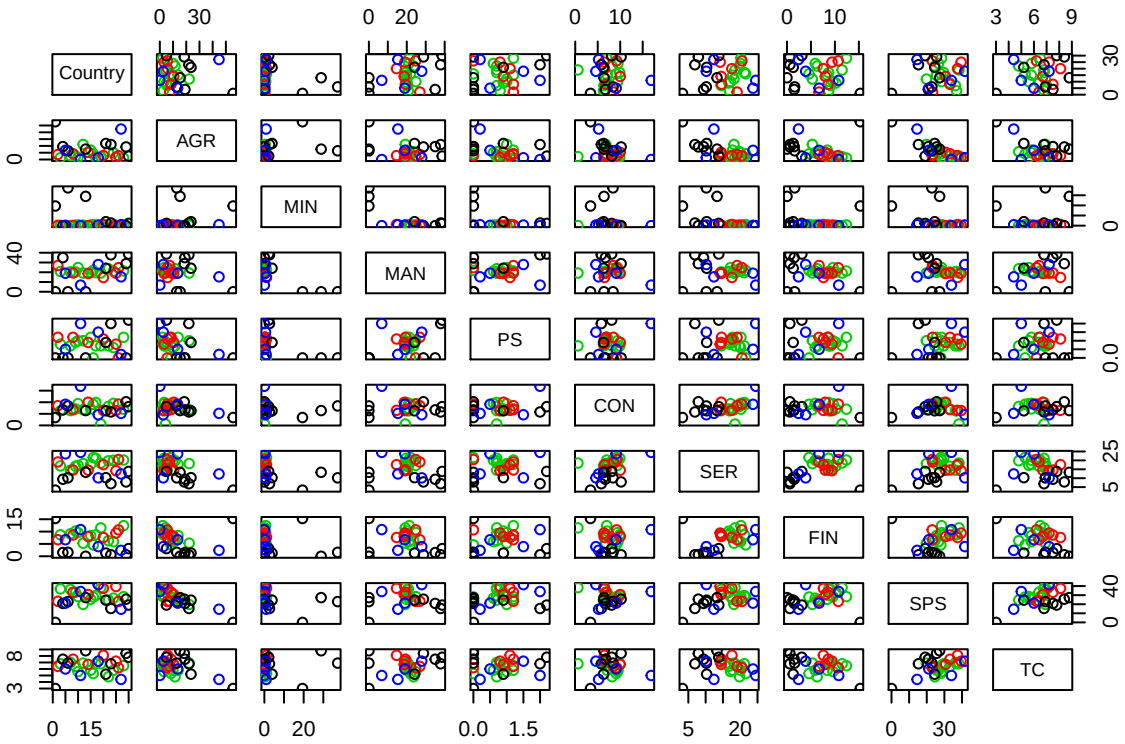

Appears that FIN vs SER and MAN vs FIN are interesting

2D plots

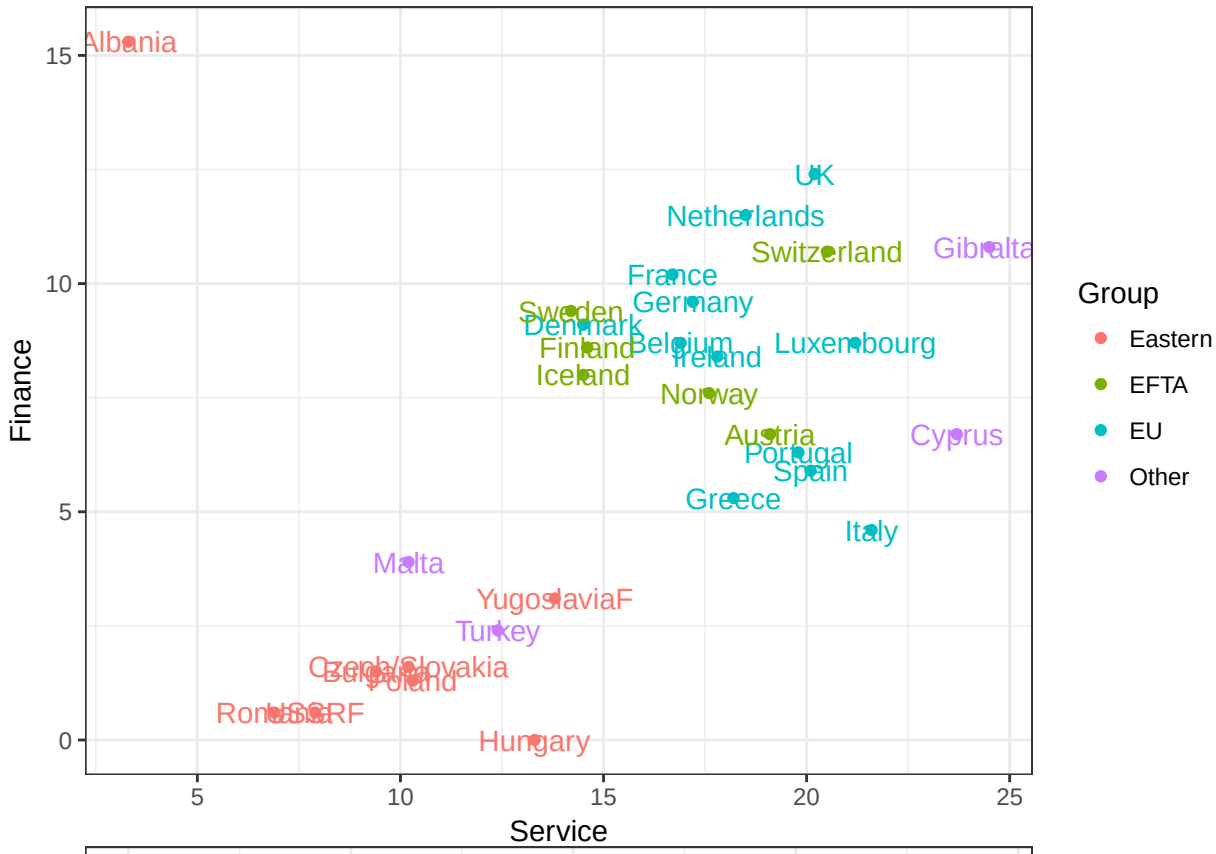
Albania
UK
Netherlands
Switzerland
Gibraltar
France
Germany
Sweden
Denmark
Finland
Belgium
Ireland
Luxembourg
Iceland
Norway
Austria
Cyprus
Portugal
Spain
Greece
Italy
Malta
YugoslaviaF
Turkey
Czech/Slovakia
Poland
Romania
USSR
Hungary
Finance
Service
Group
Eastern
EFTA
EU
Other

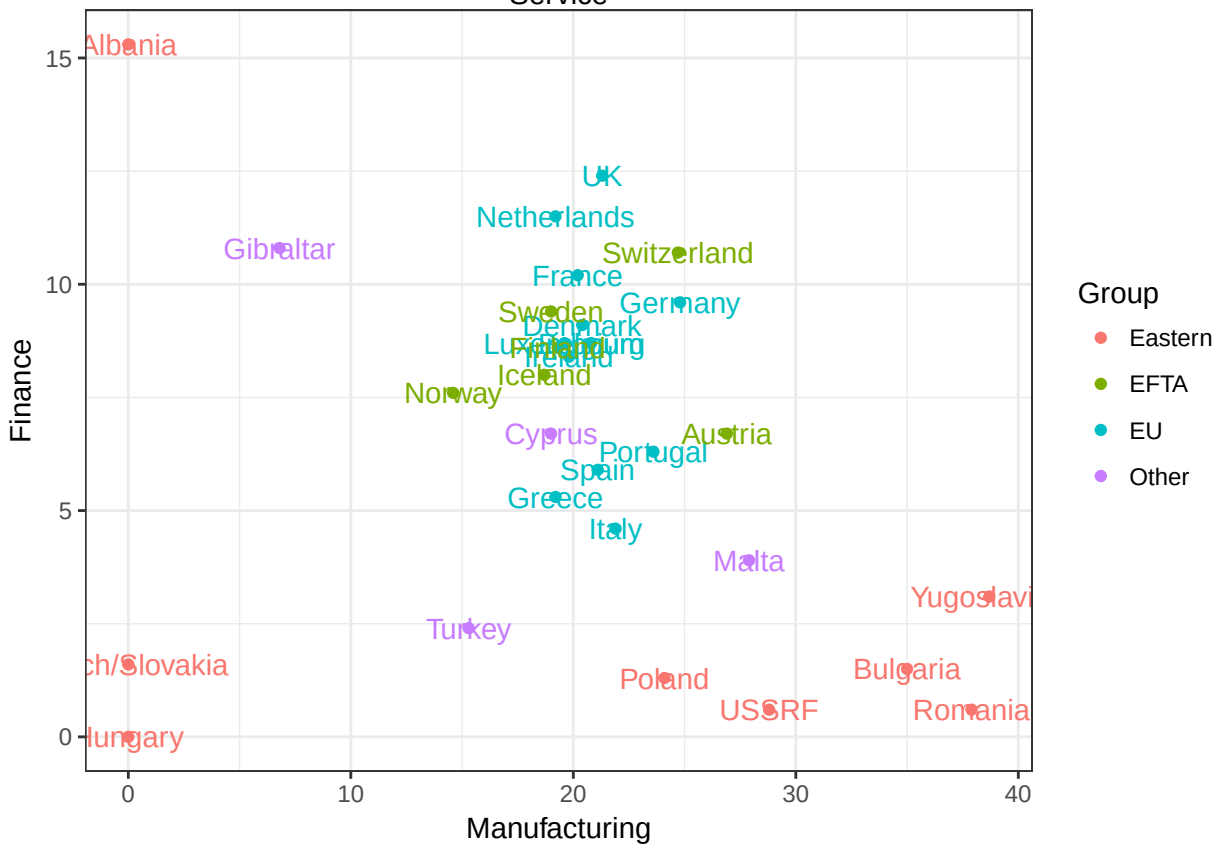
Albania
UK
Netherlands
Gibraltar
Switzerland
France
Germany
Sweden
Denmark
Iceland
Norway
Cyprus
Austria
Portugal
Spain
Greece
Italy
Malta
Turkey
Poland
Bulgaria
USSR
Romania
Finance
Manufacturing
Group
Eastern
EFTA
EU
Other

PCA

Each row is desogned to add to 100%. So if the data is standardized and scaled, we may lose the property. Hence PCA will be done without standardization.

```
pca <- princomp(my.data[, c(-1, -2)], cor = F)
summary(pca)
```

```
## Importance of components:
##                            Comp.1     Comp.2     Comp.3     Comp.4
## Standard deviation     15.380500 10.6893286 6.37294665 5.05126799
## Proportion of Variance  0.546685  0.2640566 0.09385918 0.05896537
## Cumulative Proportion   0.546685  0.8107416 0.90460079 0.96356616
##                            Comp.5     Comp.6      Comp.7       Comp.8
## Standard deviation     3.17639096 2.16103131 0.841192653 0.5459686803
## Proportion of Variance 0.02331654 0.01079241 0.001635261 0.0006888611
## Cumulative Proportion  0.98688270 0.99767510 0.999310365 0.9999992256
##                             Comp.9
## Standard deviation     1.830545e-02
## Proportion of Variance 7.743848e-07
## Cumulative Proportion  1.000000e+00
```

```
my.data$pc.1 <- pca$score[, 1]
my.data$pc.2 <- pca$score[, 2]
ggplot(data = my.data, mapping = aes(x = pc.1, y = pc.2, color = Group,
    label = Country)) + geom_point() + geom_text(size = 3, show.legend = F) +
    theme_light()
```

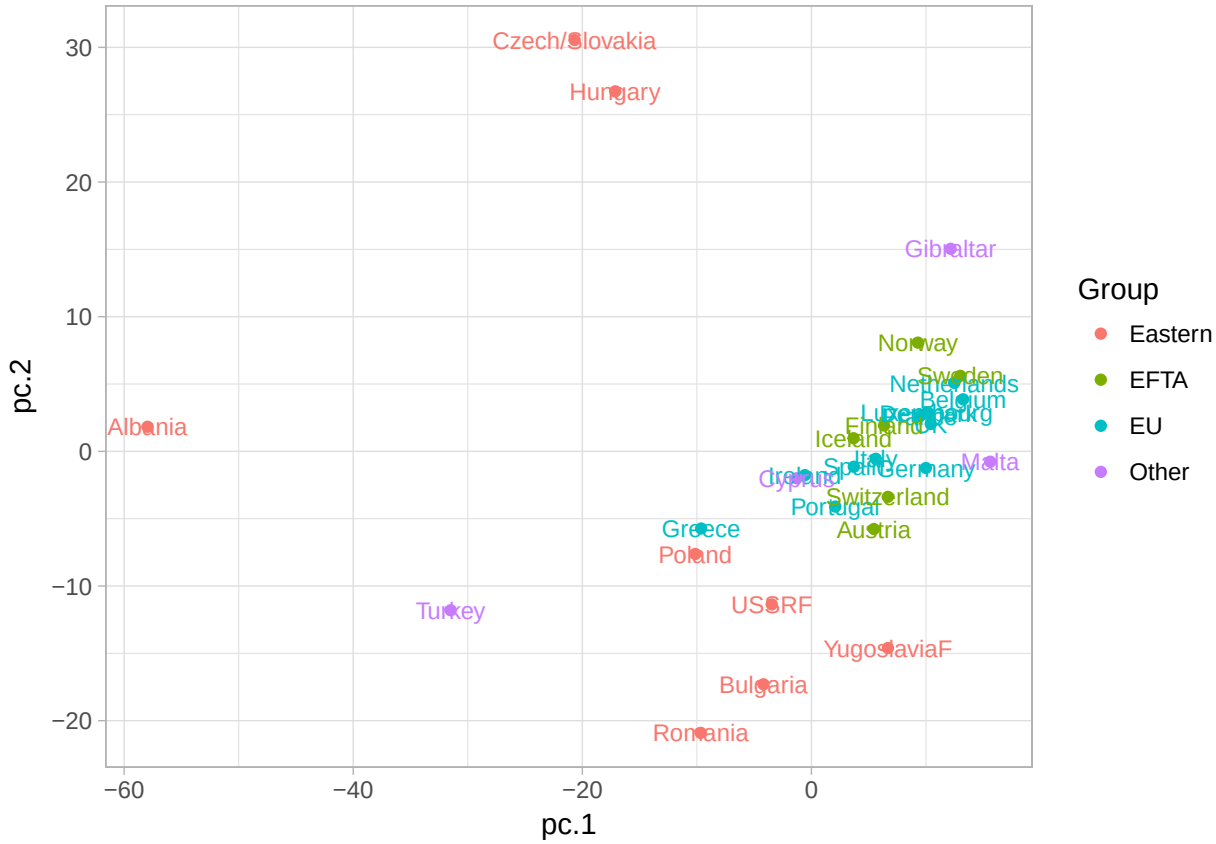

Hirearchical Clustering analysis

Let the drop the group - "other". And what does the cluster look like when we attempt to reproduce the FIN vs Ser

```
my.data <- dplyr::as_tibble(my.data)
subset <- my.data %>% dplyr::filter(Group != "Other")
fin.ser <- hclust(dist(subset[, c("FIN", "SER")]), method = "complete")
plot(fin.ser, labels = subset$Country)
```

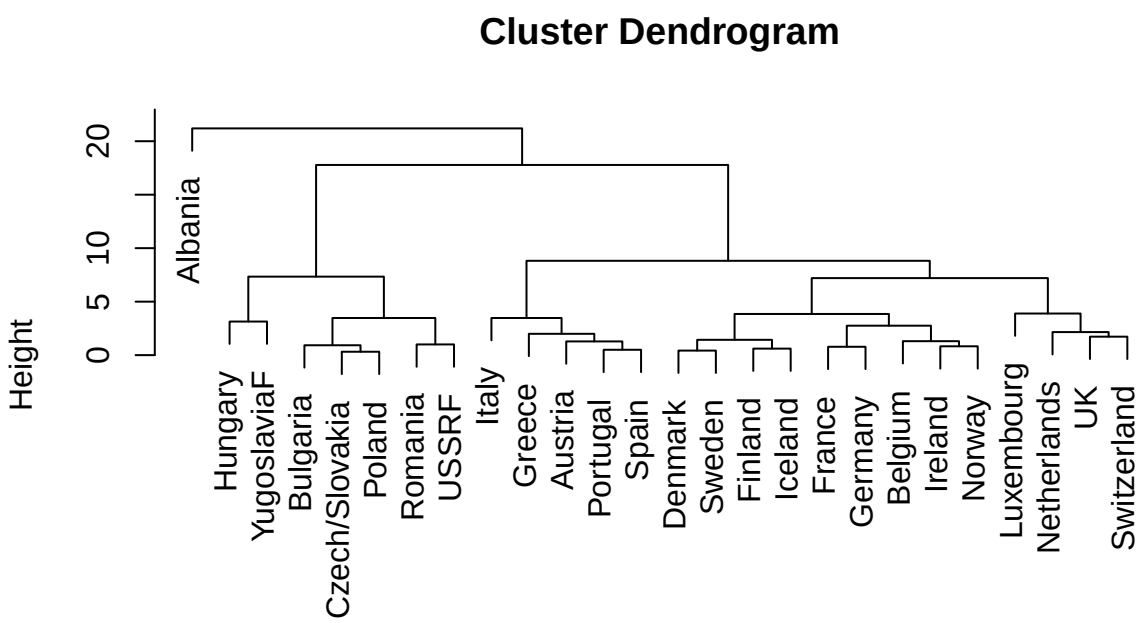


```
class <- cutree(fin.ser, k = 3)
class2 <- cutree(fin.ser, h = 12)  # class 1 & 2 are the same
subset$clust3 <- class
tab <- table(subset$Group, subset$clust3)
tab <- tab[1:3, ]
# t(tab) * (1/apply(tab,FUN = sum, MARGIN = 2)) tab

round(t(tab)/colSums(tab), 2)
```

```
##
##     Eastern EFTA   EU
##   1    0.00 0.33 0.67
##   2    1.00 0.00 0.00
##   3    1.00 0.00 0.00
```

```
class6 <- cutree(fin.ser, k = 6)
subset$clust6 <- class6
tab2 <- table(subset$Group, subset$clust6)
tab2 <- tab2[1:3, ]
```

```r
# t(tab) * (1/apply(tab,FUN = sum, MARGIN = 2)) tab
round(t(tab2)/colSums(tab2), 2)
```

```
##
##     Eastern EFTA   EU
##   1    0.00 0.44 0.56
##   2    0.00 0.20 0.80
##   3    0.00 0.25 0.75
##   4    1.00 0.00 0.00
##   5    1.00 0.00 0.00
##   6    1.00 0.00 0.00
```

Lets try a PCA again

```r
pca.out <- princomp(subset[, 3:11], cor = F)
summary(pca.out)
```

```
## Importance of components:
##                             Comp.1     Comp.2     Comp.3     Comp.4
## Standard deviation     14.9885834 10.6827988 6.06712838 4.36644154
## Proportion of Variance  0.5531982  0.2810151 0.09064127 0.04694776
## Cumulative Proportion   0.5531982  0.8342133 0.92485460 0.97180236
##                             Comp.5      Comp.6      Comp.7       Comp.8
## Standard deviation     2.72333343 1.756197858 0.848524441 0.4797367447
## Proportion of Variance 0.01826254 0.007594627 0.001772917 0.0005667161
## Cumulative Proportion  0.99006490 0.997659529 0.999432446 0.9999991616
##                               Comp.9
## Standard deviation     1.845179e-02
## Proportion of Variance 8.383717e-07
## Cumulative Proportion  1.000000e+00
```

```r
subset$pc.out.1 <- pca.out$scores[, 1]
subset$pc.out.2 <- pca.out$scores[, 2]
pca.hclust <- hclust(dist(subset[, c("pc.out.1", "pc.out.2")]), method = "complete")
plot(pca.hclust, labels = subset$Country)
```

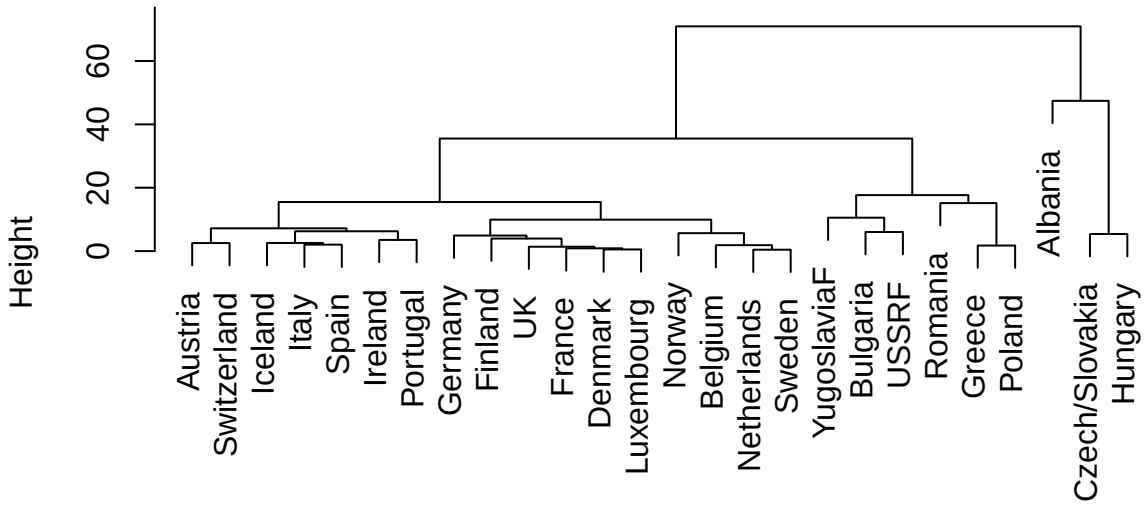


```
tab3 <- table(subset$Group, cutree(pca.hclust, k = 3))
round(t(tab3)/colSums(tab3), 2)
```

```
##
##     Eastern EFTA   EU Other
##   1    0.22 0.26 0.52  0.00
##   2    1.00 0.00 0.00  0.00
##   3    1.00 0.00 0.00  0.00
```

```
tab4 <- table(subset$Group, cutree(pca.hclust, k = 6))
round(t(tab4)/colSums(tab4), 2)
```

```
##
##     Eastern EFTA   EU Other
##   1    0.00 0.30 0.70  0.00
##   2    0.67 0.00 0.33  0.00
##   3    0.00 0.43 0.57  0.00
##   4    1.00 0.00 0.00  0.00
##   5    1.00 0.00 0.00  0.00
##   6    1.00 0.00 0.00  0.00
```

K means clustering

```
finser <- kmeans(x = subset[, c("FIN", "SER")], centers = 3)
subset$kmeans3 <- finser$cluster
centers <- as.data.frame(finser$centers)
ggplot() + geom_point(data = subset, mapping = aes(y = FIN, x = SER, color = Group)) +
    geom_text(data = subset, mapping = aes(y = FIN, x = SER, color = Group,
```

```
    label = Country), show.legend = F, size = 3, nudge_x = 0.3, nudge_y = 0.01) +
  geom_point(data = centers, mapping = aes(y = FIN, x = SER), shape = 23,
    fill = "red", size = 5) + geom_text(data = centers, mapping = aes(y = FIN,
  x = SER, label = rownames(centers)), size = 3) + theme_bw()
```

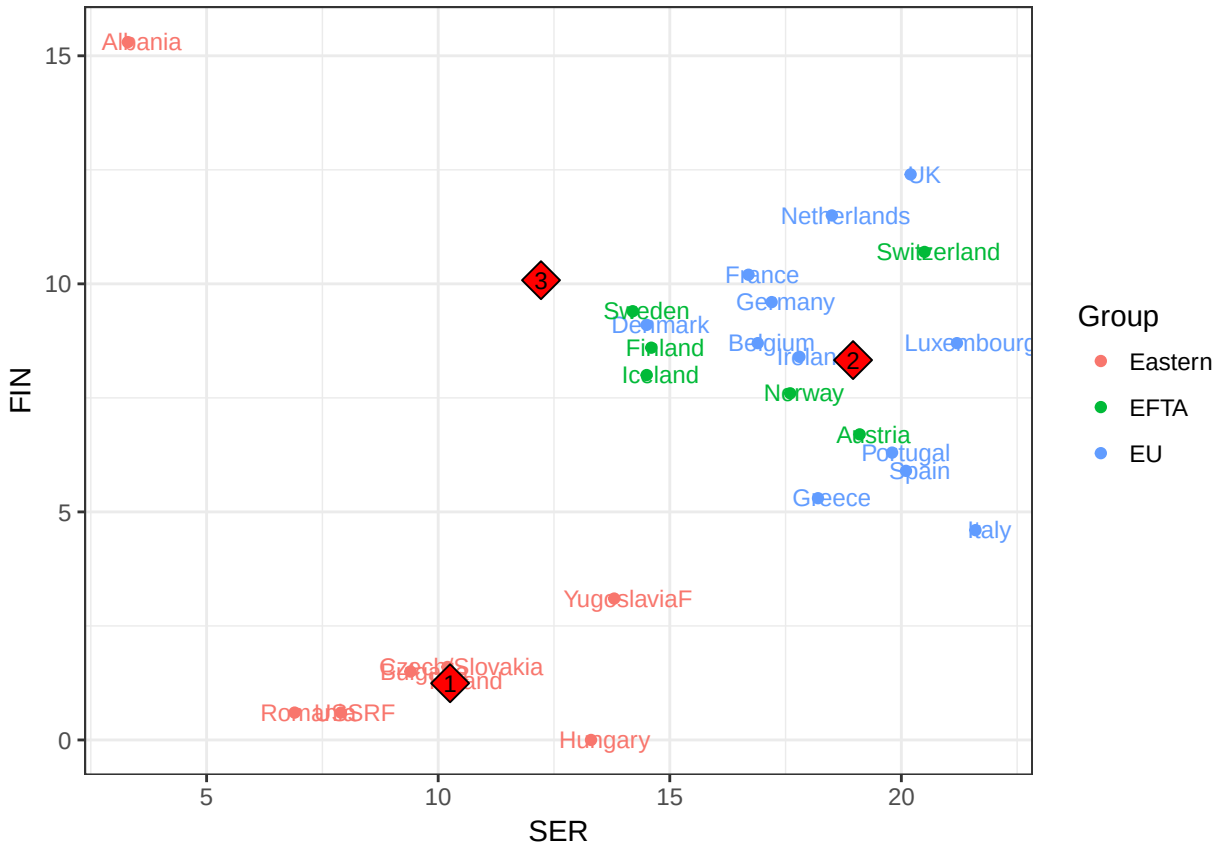


```
tab5 <- table(subset$Group, subset$kmeans3)
t(tab5)/colSums(tab5)
```

```
##
##       Eastern      EFTA        EU     Other
##   1 1.0000000 0.0000000 0.0000000 0.0000000
##   2 0.0000000 0.2142857 0.7857143 0.0000000
##   3 0.2000000 0.6000000 0.2000000 0.0000000
```

```
finser <- kmeans(x = subset[, c("FIN", "SER")], centers = 6)
subset$kmeans6 <- finser$cluster
centers <- as.data.frame(finser$centers)
ggplot() + geom_point(data = subset, mapping = aes(y = FIN, x = SER, color = Group)) +
    geom_text(data = subset, mapping = aes(y = FIN, x = SER, color = Group,
        label = Country), show.legend = F, size = 3, nudge_x = 0.3, nudge_y = 0.01) +
    geom_point(data = centers, mapping = aes(y = FIN, x = SER), shape = 23,
        fill = "red", size = 5) + geom_text(data = centers, mapping = aes(y = FIN,
    x = SER, label = rownames(centers)), size = 3) + theme_bw()
```

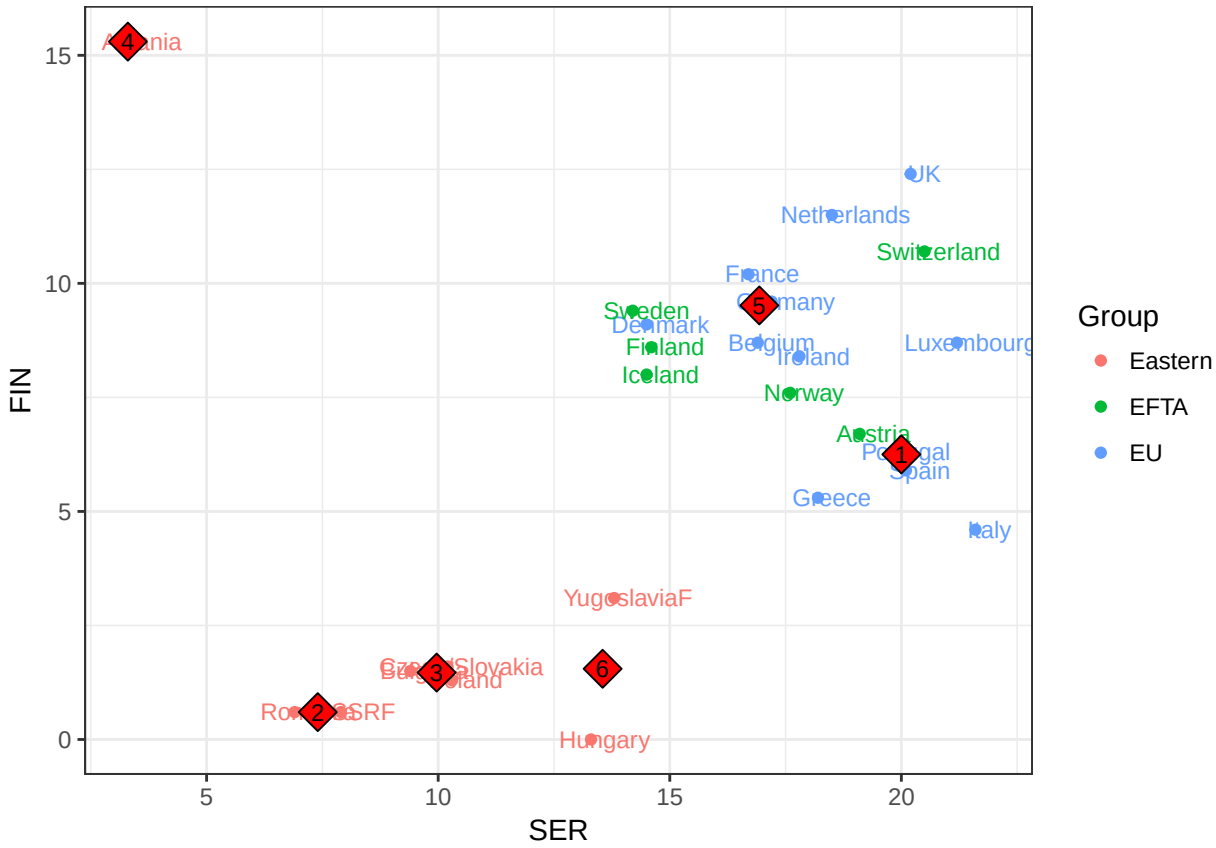


```
tab6 <- table(subset$Group, subset$kmeans6)
t(tab6)/colSums(tab6)
```

```
##
##       Eastern      EFTA        EU     Other
##   1 0.0000000 0.1666667 0.8333333 0.0000000
##   2 1.0000000 0.0000000 0.0000000 0.0000000
##   3 1.0000000 0.0000000 0.0000000 0.0000000
##   4 1.0000000 0.0000000 0.0000000 0.0000000
##   5 0.0000000 0.4166667 0.5833333 0.0000000
##   6 1.0000000 0.0000000 0.0000000 0.0000000
```

Hmmmm.... Let's do this on the principal component scores

```
finser <- kmeans(x = subset[, c("pc.out.1", "pc.out.2")], centers = 3)
subset$kmeans3 <- finser$cluster
centers <- as.data.frame(finser$centers)
ggplot() + geom_point(data = subset, mapping = aes(y = pc.out.2, x = pc.out.1,
    color = Group)) + geom_text(data = subset, mapping = aes(y = pc.out.2,
    x = pc.out.1, color = Group, label = Country), show.legend = F, size = 3,
    nudge_x = 0.3, nudge_y = 0.01) + geom_point(data = centers, mapping = aes(y = pc.out.2,
    x = pc.out.1), shape = 23, fill = "red", size = 5) + geom_text(data = centers,
    mapping = aes(y = pc.out.2, x = pc.out.1, label = rownames(centers)),
    size = 3) + theme_bw()
```

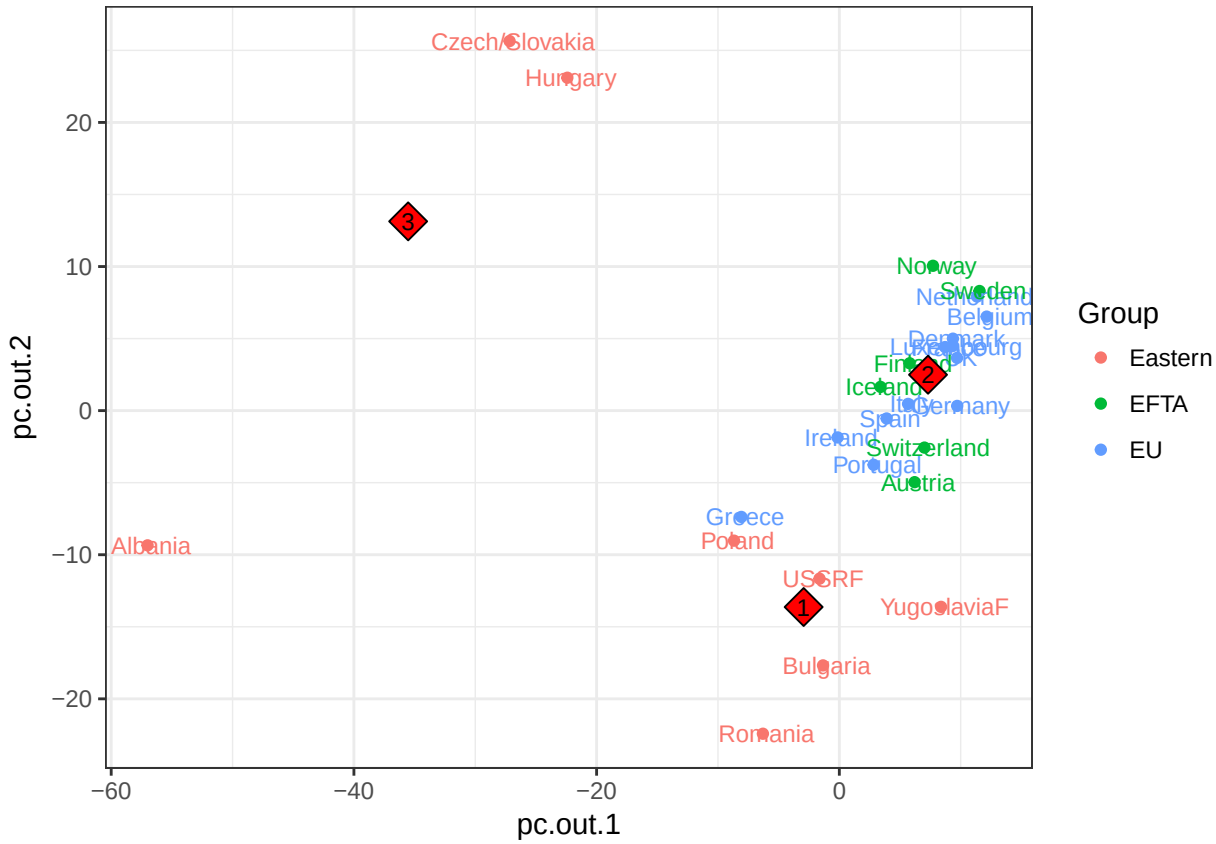


```
tab5 <- table(subset$Group, subset$kmeans3)
t(tab5)/colSums(tab5)
```

```
##
##       Eastern      EFTA        EU     Other
##   1 0.8333333 0.0000000 0.1666667 0.0000000
##   2 0.0000000 0.3529412 0.6470588 0.0000000
##   3 1.0000000 0.0000000 0.0000000 0.0000000
```

```
finser <- kmeans(x = subset[, c("pc.out.1", "pc.out.2")], centers = 6)
subset$kmeans3 <- finser$cluster
centers <- as.data.frame(finser$centers)
ggplot() + geom_point(data = subset, mapping = aes(y = pc.out.2, x = pc.out.1,
    color = Group)) + geom_text(data = subset, mapping = aes(y = pc.out.2,
    x = pc.out.1, color = Group, label = Country), show.legend = F, size = 3,
    nudge_x = 0.3, nudge_y = 0.01) + geom_point(data = centers, mapping = aes(y = pc.out.2,
    x = pc.out.1), shape = 23, fill = "red", size = 5) + geom_text(data = centers,
    mapping = aes(y = pc.out.2, x = pc.out.1, label = rownames(centers)),
    size = 3) + theme_bw()
```

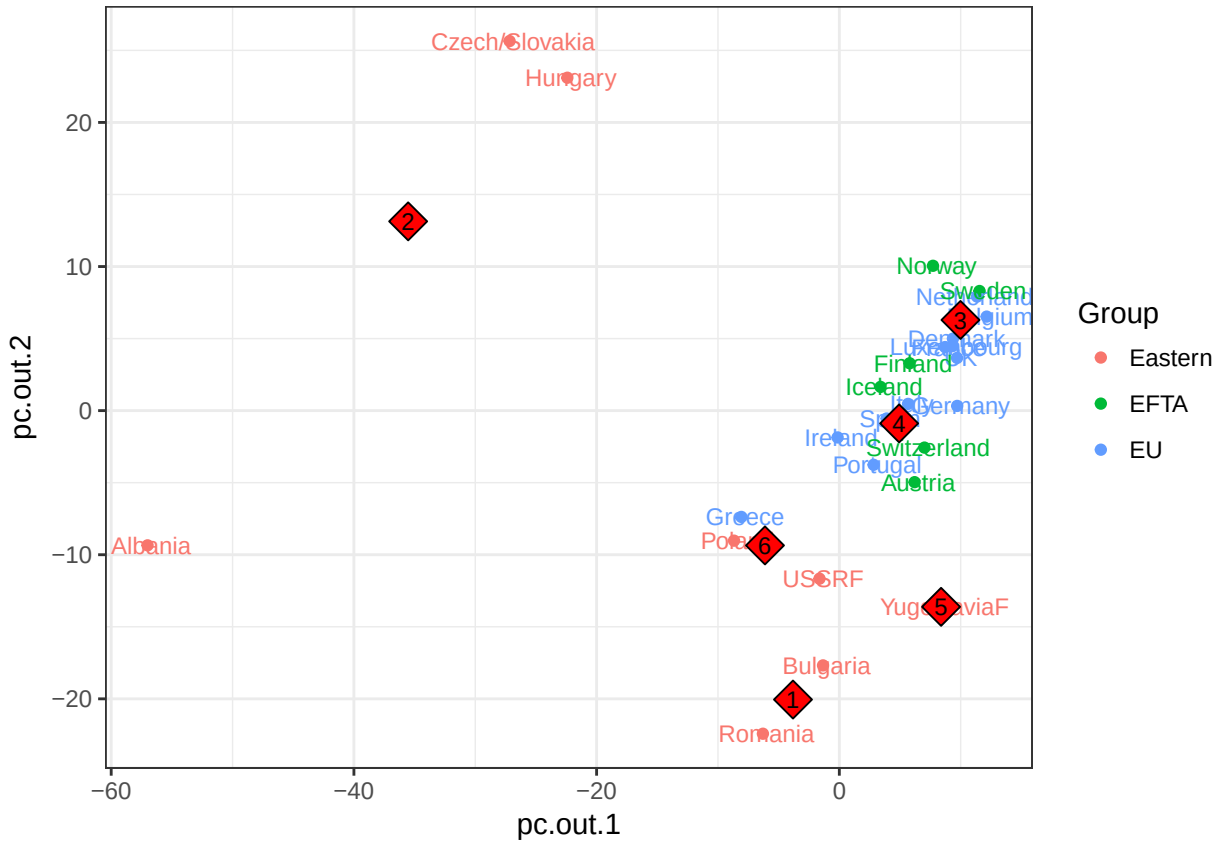


```
tab5 <- table(subset$Group, subset$kmeans3)
t(tab5)/colSums(tab5)
```

```
##
##       Eastern      EFTA        EU     Other
##   1 1.0000000 0.0000000 0.0000000 0.0000000
##   2 1.0000000 0.0000000 0.0000000 0.0000000
##   3 0.0000000 0.2500000 0.7500000 0.0000000
##   4 0.0000000 0.4444444 0.5555556 0.0000000
##   5 1.0000000 0.0000000 0.0000000 0.0000000
##   6 0.6666667 0.0000000 0.3333333 0.0000000
```

```
sum(tab5[1:3, ])
```

```
## [1] 26
```

```
Accuracy <- function(cluster, data = subset) {
    # lets get the hierarchical clustering accuracy
    hclust <- hclust(dist(data[, c("AGR", "MIN", "MAN", "PS", "CON", "SER",
        "FIN", "SPS", "TC")]), method = "complete")
    data$hiercluster <- cutree(hclust, k = cluster)
    tab <- table(data$Group, data$hiercluster)
    hier.Accuracy <- ifelse(cluster == 1, max(tab[1:3, ])/sum(tab[1:3,
        ]), sum(apply(tab[1:3, ], MARGIN = 2, FUN = max))/sum(colSums(tab[1:3,
        ])))

    # k means
    kclust <- kmeans(data[, c("AGR", "MIN", "MAN", "PS", "CON", "SER",
        "FIN", "SPS", "TC")], centers = cluster)
```

```
    data$kclust <- kclust$cluster
    tab2 <- table(data$Group, data$kclust)
    kmean.Accuracy <- ifelse(cluster == 1, max(tab[1:3, ])/sum(tab[1:3,
        ]), sum(apply(tab2[1:3, ], MARGIN = 2, FUN = max))/sum(colSums(tab2[1:3,
        ])))

    return(data.frame(cluster = cluster, hier.Accuracy = hier.Accuracy,
        kmean.Accuracy = kmean.Accuracy))
}

AccuracyResults <- purrr::map_df(.x = 1:6, .f = Accuracy)
ggplot(data = AccuracyResults, mapping = aes(x = cluster)) + geom_line(aes(y = hier.Accuracy,
    color = "red")) + geom_point(aes(y = hier.Accuracy, color = "red")) +
    geom_line(aes(y = kmean.Accuracy, color = "blue")) + geom_point(aes(y = kmean.Accuracy,
    color = "blue")) + theme_bw() + xlab("# Clusters") + ylab("Accuracy")
```